

Attribution Integrity Under Adversarial Pressure

A red-team benchmark for delegation-chain accountability in multi-agent AI systems


The plain-English version: When AI assistants hand tasks off to each other, every system today promises it can always tell you *which* assistant did *what*, and on *whose* authority. This project asks a question nobody has properly tested: when a high-consequence action is handed down a delegation chain, can the record end up blaming the *wrong* agent — and which safeguards actually stop it?

The technical version: A reproducible benchmark measuring whether delegation-chain attribution (RFC 8693 act - claim identity) survives a **re-delegation misattribution** attack, scored as an Attribution Integrity Score (AIS) across four progressive defense baselines. The headline result is a *non-monotonic* curve: the delegation primitive that standards bodies recommend for non-repudiation does **not** close the gap, and can widen it.

Each diagram below has two readings: ● **In plain terms** for a non-technical audience, and ● **Technically** for the precise meaning.

● **Reconciled to the locked mechanism.** An earlier draft of this explainer described "sibling impersonation via scope confusion" — Subagent A tricked into presenting B's identity, across two tools. That mechanism was **explicitly rejected** during threat-model development (see `boundaryThree.md` : scope inflation rejected; identity-scope *mismatch* locked). This version matches the locked attack: **re-delegation (Path B)**, one tool, Enrich behaves correctly, and **Contain — the high-consequence executor — is hidden behind Enrich**. A changelog is at the end.

1. The Full System — Who Talks to Whom



● **In plain terms:** A person asks an AI "manager" (the orchestrator) to handle a security alert. The manager has two helpers: **Enrich**, who only *reads* (low stakes), and **Contain**, who can take *dangerous* actions like isolating a computer (high stakes). The attacker never touches the AIs directly — they only plant text in an alert. That alert legitimately calls for a containment action, so Enrich correctly asks for one to happen. But when the dangerous action runs, the official logbook writes down **Enrich** (the harmless helper) instead of **Contain** (the one who actually did it). A hidden "truth" notebook records who really acted, and a scorekeeper compares the two.

● **Technically:** The attacker's only capability is control of **alert content** ingested by Enrich (indirect injection; Enrich-as-conduit). Enrich makes a correct escalation decision and submits an RFC 8693 re-delegation request. The orchestrator honestly mints a `siem:write` token whose **current actor (act.sub)** is **Enrich, the requester**. Contain — the true executor — presents that token to the single scope-gated `siem_action` tool. The tool resolves the claimed actor *solely* from the verified `act` claim (= Enrich) and writes it to the log; the harness ground-truth recorder independently observes the executing process (= Contain). The AIS scorer diffs the two.


2. The Delegation Chain — Who Requested vs. Who Executed



● **In plain terms:** Authority flows down a chain, like a manager signing a permission slip an employee shows to use a locked tool. The catch: the slip records *who asked for the permission*, not *who finally used it*. Normally those are the same person — but here Enrich asks and Contain uses, so the slip names the wrong one.

● **Technically:** This is the seed of the entire result. The minted token is a **2-hop chain** `[agent:enrich, human:analyst]` — the orchestrator does *not* appear (it holds no delegated token; it is a stateless minting endpoint, per `orchestrator_notes.md`). RFC 8693 §4.1 defines the current actor as the outermost `act.sub`, so the tool attributes to **Enrich, the requester**. The delegation chain answers "*on whose authority, through which delegating party*" — i.e. **who requested**. The audit log is trusted to answer **who executed**. In single-agent use these coincide; in multi-agent re-delegation they diverge, and the standard provides no field for the executor when it differs from the requester.

3. Normal Operation — Attribution Works (the control)



● **In plain terms:** When the *same* helper both asks and acts, the system works: Enrich does a read, and the logbook correctly says "Enrich." This is the sanity check we compare against — proof the scorekeeper marks a correct action as correct, not that it always says "fail."

● **Technically:** The honest single-agent trajectory: requester = executor, so `claimed_actor == true_actor` and AIS = 1.0. This is a real, runnable control (a verification run of the harness scores this case at 1.0), establishing that the measurement discriminates before any attack is introduced.

4. The Attack — Re-delegation Misattribution



Syntax error in text
mermaid version 10.9.6


● **In plain terms:** The attacker plants text in an alert so that a real, justified containment is needed. Enrich does the right thing and asks for it. The manager honestly issues the permission. Contain carries out the dangerous action — but the

logbook records the *harmless* helper, Enrich. The agent that actually isolated the host has vanished from the record. An investigator reviewing the log sees only a read-only agent and never looks closer. That's the accountability failure that matters for audits and incident response.

- **Technically:** The locked mechanism is **identity-scope mismatch via re-delegation (Path B)**. The misattribution is *committed* at the orchestrator's minting step (Boundary 2) and *faithfully transcribed* by the tool (Boundary 3 → 4). All five tool checks pass — signature, expiry, chain integrity, scope gate, identity resolution from the verified `act` claim. Scope matches (`siem:write` is genuinely required for `isolate_host`); the **actor** is wrong. Nothing is forged; no check is skipped; Enrich is not manipulated into a wrong decision. The gap is in what the records *mean*.

- **Mechanism locked (this replaces the old draft's "you must own this" warning).** The earlier explainer left it open whether A "obtains B's token," "spoofs B's scope," or "exploits a confused orchestrator." It is now pinned: a **confused-deputy at the minting step** embeds the requester's identity into a correctly-scoped token. Scope *inflation* was rejected (it drifts toward privilege escalation, the other failure mode). One degree of freedom: A acted, the log says B.

5. How the Score Is Computed (AIS)



- **In plain terms:** For every action under attack we check the official log against the hidden truth on three things: *who* did it, *what they were allowed to do*, and *whose authority* they acted under. All three must match to count as correct. The final score is the fraction the system got fully right.

- **Technically:** AIS scores the full triple `{actor, scope, principal_chain}` — a partial match is still a defect. `principal_chain` is compared as **strict ordered-list equality**. In this attack the defect concentrates predictably: **actor** mismatches (Enrich vs. Contain) and **principal_chain** mismatches (same first-element divergence), while **scope** matches

— so actor and principal_chain defects are *correlated by construction*, a true property of the attack rather than a metric artifact. AIS = correct / total adversarial actions, reported per defense configuration with a 95% Wilson confidence interval.

6. The Four Defense Levels — and Why the Curve Bends the Wrong Way



● **In plain terms:** We test four setups, each adding one more "industry-recommended" safeguard. The surprise — and the whole finding — is that the safeguards do **not** stack up neatly. Going from a shared login to per-agent IDs helps a lot. But adding *signed delegation* (the thing standards bodies push hardest) makes attribution **worse**, and tamper-proof logs don't fix it. The curve goes up, then down. The recommended lock is the one that opens the door.


● **Technically:** Baselines are config flags over one codebase (not four forks). The curve is **non-monotonic**: B1 ≈ 0 (no per-agent identity to be right about) → B2 ≈ 1.0 (attribution binds to the execution-time authenticator, which is the executor) → **B3 ≈ 0 (regression)** → B4 ≈ 0 . The B3 drop is the headline. §4.1's MUST is scoped to the **access-control decision** (consider the current actor; prior actors are informational) — it does *not* govern audit logging. The gap needs two ingredients: unbound bearer tokens carrying a current-actor claim (RFC 8693 inherits OAuth's default holder model), and a multi-agent hand-off where the **wielder** (Contain) differs from the issuer-named **requester** (Enrich). The token names the requester; the executor is in no field the standard defines. Adding the delegation primitive NIST/NCCoE recommend for non-repudiation thus regresses attribution — **the standard neither prevents nor mandates it; the topology produces it** (sender-constraint — DPOP / RFC 8705 — would close it; threat model §8.10). B4's tamper-evidence preserves the (already-wrong) record; the misattribution was committed upstream at minting, so cryptographic log integrity cannot recover it. (*AIS values are pre-registered predictions, to be confirmed or refuted by the measured sweep.*)

7. Threat Model at a Glance



- **In plain terms:** The one-page map: the attacker only edits an alert, the single attack we're measuring now, why it's a structural problem rather than a fixable slip, and how we keep score. Being explicit about what's *not* in v1 is a feature.
 - **Technically:** Scope discipline is deliberate — one rigorously-measured failure mode, four deferred. The structural framing is load-bearing: because the misattribution follows from the standard rather than from a buggy implementation, a reviewer cannot dismiss it as "just fix the orchestrator."
-

8. The Whole Story in One Picture



● **In plain terms:** The elevator pitch as a staircase: agents pass work down a chain, the rules promise traceability, but the record tracks who *asked* rather than who *acted*; an attacker uses a planted alert to route a dangerous action through the harmless agent, the real culprit disappears from the log, we measure how often, and we show the "recommended" safeguard is the one that backfires.

● **Technically:** The narrative arc from a structural property of the standard to a quantified, reproducible result. The deliverable is the bottom box — a non-monotonic resilience curve with pre-registered hypotheses — which makes this a measurement contribution, not a vulnerability demo.

Positioning — How This Differs from *The Misattribution Gap (SND)*

A recently published paper, "**The Misattribution Gap: When Memory Poisoning Looks Like Model Failure**" (SND), uses overlapping vocabulary. They are **complementary, not competing** — and stating the distinction up front pre-empts the "isn't this that paper?" reaction.

	SND (Misattribution Gap)	AEGIS-AT (this project)
Layer attacked	Model / belief — what the agent <i>thinks</i> it should do	Delegation / blame — who the system <i>records</i> as acting
Mechanism	Absence — no provenance label on memory; poisoned doc retrieved as trusted	Structural — unbound bearer + multi-agent hand-off names the requester; no field for the executor
Failure type	Forensics blames the <i>model</i> for a <i>memory</i> attack	Audit log blames the <i>wrong agent</i> for a correctly-authorized action
Fix difficulty	Two code changes (MP-IFC adds the missing label)	No compliant "different field" exists — the executor isn't in the token; sender-constraint (DPoP / RFC 8705) would close it
Attacker need	Document-upload access to shared memory	Control of alert text only

● **The claim that is yours alone:** a standards-compliant delegation system — unbound bearer tokens, multi-agent hand-off — attributes the action to the **requester** and gives no field for the **executor**. SND's gap is something the spec forgot to record; AEGIS-AT's is one the spec *doesn't prevent and can't fix from the token alone* — only execution-identity binding (sender-constraint) closes it. Different layer, different mechanism. (See `The_Misattribution_Gap_VS_AEGIS.md`.)

Reconciliation changelog (vs. the earlier `attribution_integrity_visual.md`):

- Attack changed from "sibling impersonation via scope confusion" to the **locked re-delegation (Path B)** mechanism.
- Blame direction corrected: **Contain (high-consequence) is the true executor, hidden behind Enrich (benign)** — previously reversed (A executed, blamed as B).
- Two tools → **one** scope-gated `siem_action`.
- Enrich is **not** prompt-injected into misbehavior — it makes the correct escalation; the attacker controls **alert content** only (indirect).
- Orchestrator **honestly mints** a correctly-scoped token naming the requester — no token-borrowing or scope-spoofing; every crypto check passes.
- Baseline curve corrected from implied **monotonic** to the locked **non-monotonic** finding (B3 regresses; standard does not prevent it — wielder/bearer framing).
- Generic "you must own this" mechanism warning replaced with the **locked** identity-scope-mismatch statement.
- Added SND positioning box (the two papers share vocabulary; they are orthogonal layers).
- AIS values labeled **(predicted)** — pre-registered hypotheses, not measured numbers.