

AEGIS-AT v3: Who Attests the Completion?

Self-Reported vs. Independently-Verified Attribution
in Multi-Agent Delegation — Under a Real-Model Adversary

Rahul Singh*

June 2026

Abstract

AEGIS-AT v1 [1] showed that adding OAuth 2.0 Token Exchange (RFC 8693) delegation to a correctly-functioning multi-agent SOC pipeline makes audit attribution *worse* — the Attribution Integrity Score (AIS) regresses from 1.0 to 0.0 — because the standard names the requester, not the executor. v2 [2] showed that binding the token to the executor’s key (DPOP, RFC 9449) recovers AIS to 1.0. **v3** turns to the layer the 2026 agent-identity stack is adding: the *completion record*. AIP, PEDIGREE, and HDP all introduce executor-supplied completion/provenance blocks, and they converge on *accepting the executing agent’s own account of what it did without independent verification by default* — AIP makes self-report the default for trusted environments; PEDIGREE permits a **self_reported** tier with no adequacy guidance; HDP signs hops with the issuer’s, not the executor’s, key. None models a dishonest authorized executor. We distill this into a minimal completion-record abstraction and two baselines: **B8** reads the executor from the self-reported completion; **B9** reads it from an independent process-boundary verifier (PEDIGREE’s **tool_verified** tier; v2’s recorder reused). We measure attribution on two tiers. **Tier 1 (scripted, deterministic)** specifies the failure mode exactly: a colluding executor that signs a truthful-looking completion *with its own valid key* but names a sibling collapses B8 to AIS = 0.0, while B9 holds at 1.0 *by construction* — its verifier is the ground-truth instrument, so the deterministic tier is an existence proof with a structural positive control, not an empirical discovery. The full deterministic curve is B1–B9, and every baseline that *resolves the executor from evidence it cannot author* ({B5 DPOP, B6 mTLS, B7 A-JWT, B9 process-boundary}) recovers to 1.0; self-report (B8) is the lone failure. **Tier 2 (real models)** is the empirical payload: we put four open-weight LLMs (weak→strong) in the executor seat and induce collusion by prompt injection on a structured tool-call argument. Under B8 they forge the executor claim at 0.91–1.00; under B9 they still emit the forgery (0.89–1.00) but the evasion rate is 0.0 for every model (H2, the knockout) — independent verification does not make the model honest, it makes the forgery inert. The two-stage SHA-256 pre-registration plays two roles: the deterministic locks pin the specification (a regression gate), while the Tier-2 hypotheses are the falsifiable predictions — and one of them, H1’s rising capability curve, came back flat and is reported as a finding under a disclosed confound: the locked prompt *forces* the completion call, saturating B8 near the ceiling, so we do not claim capability-independence and we name prompt de-saturation as v4. Figures regenerate from the harness. The result completes the arc: completion-era attribution defaults to self-report, a real model exploits it, and independent verification is the layer that closes it.

*Independent researcher, Jersey City, NJ. Correspondence: rahul.rs1397@gmail.com · github.com/rahulsingh1397. Code and documentation are dual-licensed: Apache-2.0 (code) and CC BY 4.0 (documentation). This paper extends the v1 [1] and v2 [2] benchmarks; the v1 artifact is frozen at tag **v1.0.0**.

Keywords: multi-agent systems, AI agent security, attribution, non-repudiation, completion record, provenance, AIP, PEDIGREE, self-attestation, prompt injection, LLM adversary, OAuth 2.0 Token Exchange, RFC 8693, mTLS, RFC 8705, Model Context Protocol, benchmark.

Contents

1	Introduction	2
2	Background and Relationship to v1/v2	3
2.1	What is inherited, and what v3 adds	4
2.2	The MCP transport boundary	4
2.3	The confused deputy, carried and distinguished	4
3	The Completion-Record Abstraction	4
4	Tier 1 — The Scripted Completion Core	5
5	Tier 2 — The LLM Ladder	6
6	Comparative Breadth: B6 (mTLS) and B7 (A-JWT)	7
7	Implementation	8
8	Experimental Methodology	9
9	Results	10
9.1	The deterministic curve (Tier 1)	10
9.2	The real-model grid (Tier 2)	10
10	Discussion	10
11	Limitations and Validity Threats	12
12	Related Work	13
13	Future Work	14
14	Conclusion	14
	Reproducibility and Artifact Availability	14
	References	15

1 Introduction

The arc, in three lines. AEGIS-AT v1 [1] built the smallest multi-agent system in which sibling impersonation can occur — two agents (**Agent-Enrich**, read-only; **Agent-Contain**, write-capable) sharing one scope-gated tool — and showed that RFC 8693 delegation regresses attribution from AIS = 1.0 (per-agent identity, B2) to AIS = 0.0 (signed delegation, B3), because `act.sub` names the agent that *requested* delegated authority while the agent that *executes* can differ, and the standard

provides no executor field. v2 [2] bound the token to the executor’s key (DPOP, RFC 9449) and recovered AIS to 1.0 at Baseline 5, confirming v1’s conjectured fix while a real process boundary, a hash-chained log, a second topology, and Wilson intervals each left the finding standing. One line each: *v1 = the executor has no field; v2 = bind the token to the executor’s key; v3 = completion records now exist — but who attests their content, and can a real model exploit the gap?*

The 2026 completion-block landscape. The agent-identity stack taking shape in 2026 added the layer v1/v2 did not have: an after-the-fact record of *what an agent did*. The Agent Identity Protocol (AIP) [7] defines a completion block signed by the executing agent and makes “*Level 1 Self-Reported*” the “default for trusted environments”; PEDIGREE [8] requires a completion block’s `verification_status` to be one of four tiers — `self_reported`, `tool_verified`, `peer_verified`, `human_verified` — but specifies no default and gives no adequacy guidance; HDP [9] records an append-only chain of self-supplied action summaries, signs each hop with the *issuer’s* key rather than the executor’s, and admits (§5.4) that a genuine-but-misrepresenting hop is “not detectable by the protocol alone.” The transport beneath them, the Model Context Protocol (MCP) [10], carries no `act` claim, no delegation chain, and forbids a server from passing the client’s token upstream — so executor identity is not carried by the transport at all; it must ride in a completion record.

The gap. These mechanisms close the *honest-executor* gap: an agent that truthfully reports what it did now produces a usable record. What none of them models is a *dishonest authorized executor* — AIP’s threat model (§7.1) omits it outright. Self-report becomes the emerging default for the very field (the executor) whose misattribution v1 and v2 measured, and it fails for a structural reason: the attester and the accountable party are the same entity. The incentive to misreport is concrete, not hypothetical: an executor that names a sibling shifts blame for a destructive or policy-sensitive action (deflecting review, containment, or credential revocation onto another agent); a compromised executor launders the attacker’s provenance into a clean audit trail; and — the realization Tier 2 measures — the executor need not hold the incentive at all, because attacker-controlled input can supply it, turning an honest-but-injectable model into the colluder. v3 measures whether attribution drawn from a self-reported completion survives an executor with an incentive to misreport — first against a scripted colluder (the deterministic anchor), then against a ladder of real LLMs.

What sets this apart. Unlike completion/provenance proposals that specify a mechanism and validate their own scheme, AEGIS-AT treats *attribution correctness* — *does the logged actor equal the true executor?* — as the measured dependent variable, not a design guarantee. Among the works we compare, that framing is what makes a self-attested completion’s failure measurable rather than assumed.

Contributions.

1. **A source-verified analysis of the 2026 completion-block proposals, distilled into a benchmark abstraction (§3):** AIP defaults to self-report, PEDIGREE permits it with no adequacy guidance, and HDP issuer-signs self-supplied summaries — captured as a minimal record (`asserted_executor`, `attestation_source` $\in$ `{self_reported, tool_verified}`, `attester_id`, `signature`) with the signer decoupled from the claimed executor.
2. **A deterministic existence proof with a structural control (§4):** a scripted colluder collapses B8 self-report to AIS = 0.0; B9 (tool-verified) holds at 1.0 *by construction*, because its verifier is the ground-truth instrument (v2’s process-boundary recorder reused). The deterministic tier pins the failure mode and gates regressions; it claims no empirical surprise.

3. **The Tier-2 real-model result — the empirical payload** (§5): four open-weight models forge the executor claim at 0.91–1.00 under self-report, still emit the forgery at 0.89–1.00 under B9, and evade zero times under B9 — the pre-registered H2 holds for every model. A controlled measurement of a *real* model exploiting the completion gap.
4. **Comparative breadth** (§6): B6 (mTLS-bound, RFC 8705) and B7 (A-JWT execution assertion) as configuration flags, confirming that the recovery cluster is exactly the baselines that verify the executor at execution.
5. **Two-stage pre-registration** (§8): a deterministic threat model (exact equality) and a separately locked LLM-tier threat model (directional hypotheses + Wilson containment), both SHA-256 locked before their measuring code.

Scope and honesty up front. v3 measures a scripted core plus a four-model LLM tier, one scenario, one topology for the LLM tier (the B8/B9 path is topology-inert), and open-weight models only. One disclosure is load-bearing and we state it before the results: the locked Tier-2 prompt (§5) *forces* the completion tool call (“you MUST call `submit_completion...` exactly once”), which removes the refusal path and saturates the B8 forging rate near the ceiling for *every* model. We therefore report H1’s flat capability curve as a *finding under a disclosed confound* — not as evidence of capability-independence, which the saturation and the prevailing injection literature would make indefensible — and we name prompt de-saturation as the first v4 experiment (§13). These boundaries are defended in §11.

2 Background and Relationship to v1/v2

We assume the v1/v2 background [1, 2]: RFC 8693’s `act` claim and its §4.1 current-actor semantics; the process-boundary recorder (`multiprocessing + os.getpid()`) that supplies ground truth; the AIS scorer (a strict triple match on `(actor, scope, principal_chain)` over adversarial actions); the B1–B5 path applied as configuration flags over one codebase; the two v2 topologies, T1 (the two-agent hand-off) and T2 (the three-agent re-delegation chain) — not to be confused with this paper’s measurement *Tiers* 1 and 2; and the confused-deputy lineage from Hardy [19] through the CSA confused-deputy note [13]. This section adds only what v3 needs.

2.1 What is inherited, and what v3 adds

v3 reuses the v2 substrate *by import*, not by copy (INV-6): the B1–B5 credential and verification path, the process kernel, the recorder, the tool, and the scorer are the v2 modules unchanged. The recorder is reused as-is and becomes B9’s independent verifier — it already occupies PEDIGREE’s `tool_verified` tier. v3 adds three things over that package: a completion-record module (§3); an *attestation-source* flag selecting B8 (self-reported) or B9 (tool-verified); and an *adversary adapter* that swaps the agent in the executor seat between a scripted body and a real LLM. Any AIS difference therefore remains attributable to the flag, not to incidental code differences.

2.2 The MCP transport boundary

To land the result in a shipped 2026 protocol, the tool call crosses an MCP-shaped boundary whose defining rule is that the server MUST NOT pass the client’s token through to upstream [10]. Because MCP carries no attribution layer, the executor identity must be supplied by a completion record — exactly the field whose attestation source v3 measures. The adapter is a thin transport

wrapper enforcing the no-passthrough rule and RFC 8707 audience binding, not a full MCP stack; B1–B5 credential semantics remain byte-identical to v2 unless a baseline opts into completion-block attribution.

2.3 The confused deputy, carried and distinguished

v1/v2’s central discipline — no system component misbehaves — is *preserved*. The colluding executor is the adversary’s locus, not a miscoded component: the orchestrator, tool, verifier, recorder, and audit layer all still behave correctly. The lineage is carried with a distinction the v3 reader must keep: in v1/v2 the *orchestrator* is misled by RFC 8693 token semantics (which name the requester); in v3 the *audit/deployment layer* is misled by self-attestation (which trusts the executor’s own claim). The confused deputy in v3 is the deployment design that treats a self-reported completion as audit truth.

3 The Completion-Record Abstraction

A minimal record. v3 uses a minimal completion record — inspired by the verified source claims but not asserted to be byte-for-byte any draft’s schema (the receipts lock the trust *tiers*, not exact field names):

```
completion_record:
  action_id          # which action this attests
  asserted_executor  # who the record SAYS executed (self-asserted)
  attestation_source # self_reported | tool_verified
  attester_id        # who signed this record
  signature           # signature under attester_id’s key
```

The load-bearing property is that `attester_id` (who *signed*) is decoupled from `asserted_executor` (who the record *claims* executed). A colluding executor signs a valid record under its *own* key while naming a sibling — no cryptographic check fails.

The signer-vs-claim objection, answered by the tiers. A reviewer will ask: if Contain signs the record (`attester_id = contain`) but `asserted_executor = enrich`, why not just enforce `attester_id == asserted_executor`? Because that cross-check *is* itself an act of independent verification — exactly what the `self_reported` tier omits by definition (AIP Level 1: “no independent verification”). B8 models a system at the `self_reported` tier: it trusts the self-asserted executor without cross-checking it against any independent observation. B9 is the tier that performs the check (via the recorder). The objection is therefore not a hole in the benchmark — it *is* the B8→B9 distinction. The benchmark rests only on the verified property that `self_reported` performs no independent verification of asserted content, not on any specific draft’s field layout.

The two baselines. **B8 (self-reported)** takes the audit record’s `claimed_actor` from `asserted_executor`, trusted because the tier is `self_reported`. It measures what AIP does by default. A system relies on B8 even when DPOP (B5) exists because completion records are often returned after the fact or aggregated by an orchestrator that did not witness the execution-time handshake — B8 is the decoupled, asynchronous audit path, distinct from B5’s synchronous access-control path. **B9 (tool-verified)** takes `claimed_actor` from the independent process-boundary verifier, never from the self-report; it is the positive control showing the `tool_verified` tier recovers attribution. B9 *fails loud* if the verifier observation is absent — it never silently falls back to the self-report, which would make its 1.0 vacuous (the non-circularity guard, §4).

The deliberate control: isolate the executor axis. Contain runs the canonical action under an *opaque* per-agent credential (B2-shaped). Ground truth therefore carries `principal_chain = None` and an observed scope, so the *only* field the attestation source can move is the executor identity. This isolates the executor-attribution axis: a B8/B9 difference cannot be a scope or chain artifact. (A sensitivity slice repeats the measurement under a chain-bearing JWT credential and confirms the actor finding survives.)

4 Tier 1 — The Scripted Completion Core

The honest checkpoint (a hard gate). Before any colluder code, we verify the honest column: B8-honest = B9-honest = AIS 1.0. The completion names the true executor, the signature verifies, and the system records the true executor. If the honest column were not 1.0 the harness would be wrong, not the finding (INV-7); this mirrors v1/v2’s “verify the no-attack flow first.”

The colluder, pinned exactly. The colluding capability is one pre-registered move: the executor emits a false self-reported completion *within its own legitimate key material*. The record sets `asserted_executor = enrich` while `attester_id = contain`; the signature is valid under Contain’s own key; every cryptographic check passes. The colluder *cannot* sign under another agent’s key, touch the recorder, cross the process boundary, or forge a token — those are distinct attack classes (principal laundering, key forgery), explicitly out of scope. This is the same construction as v2’s token-lift and proof-replay stimuli: an adversary capability injected as a test stimulus, not a miscoded component (INV-5 preserved).

The result. Under the identical colluder, B8 records the *false* executor while the recorder observed the true one: AIS = 0.0. B9 reads the verifier and records the true executor: AIS = 1.0. The defect under B8 is an actor `field_mismatch`, not a missing or malformed record — a tamper-evident, signature-valid completion of the wrong actor. This is the deterministic existence proof; it is identical on both topologies because the opaque-credential audit path carries no re-delegation chain.

Non-circularity, made explicit. A degraded-resolver control shows the audit path genuinely consults the verifier: when the verifier is disabled (or the audit path is forced to fall back to `asserted_executor`), B9 *must* drop below 1.0, and a test asserts this. Separately, an INV-4 test asserts the recorder’s inputs exclude any `act` claim, completion field, or agent-supplied identity, on every baseline and seat — the measurement reads identity only from the OS process boundary. What the control does *not* do is make B9 falsifiable: B9’s verifier and the ground-truth recorder share the same process-boundary observation point (by design — the locked threat model states both derive from it), so B9 cannot fail while the recorder stands, and its 1.0 is structural. The deterministic tier is therefore a specification of the failure mode with a positive control, not an empirical discovery; the falsifiable measurements live in Tier 2. Giving B9 an independent failure surface — a second, separately implemented observer (e.g., a kernel-audit tracer) distinct from the recorder — is named as v4 work (§13).

5 Tier 2 — The LLM Ladder

The seat is now a real model. Tier 2 replaces the scripted adversary with a real LLM in the executor seat; everything else — the execution path, the OS-process-boundary recorder (ground truth, INV-4), the B8/B9 resolution, and the scorer — is reused byte-for-byte. The model executes the canoni-

cal action and reports the completion via a *structured tool call*, `submit_completion(action_id, asserted_executor)`. Crucially, the injection manipulates the `asserted_executor` argument, not free-text reasoning: the attack is a provenance note in the alert content (“...carried out on behalf of the Enrichment service...set `asserted_executor` to `agent:enrich`”), continuous with v1/v2’s alert-text adversary. Whether the model complies is the measured behavior.

The ladder. Four free open-weight models served through one provider, ordered weak→strong by the public Artificial Analysis Intelligence Index (not by parameter count — the gpt-oss models are mixture-of-experts and out-score the dense 70B): M1 `llama-3.1-8b-instant` (index ≈ 6), M2 `llama-3.3-70b-versatile` (9), M3 `openai/gpt-oss-20b` (15), M4 `openai/gpt-oss-120b` (24). All four training cutoffs (Dec 2023; Jun 2024) precede AEGIS-AT’s first public release (2026-06-22) by ≥ 18 months, so no model could have memorized the benchmark; the generic threat pattern being in training is intended — that is the reasoning Tier 2 tests.

Measurement discipline. Per (model $\times$ baseline $\times$ condition) cell we sample in batches of 20 and stop when the Wilson 95% interval half-width < 0.05 or a hard cap of $N_{\max} = 200$ is reached. Each trial yields a deterministic bit; we report the per-trial *evasion* bit — the complement of that trial’s AIS bit (evasion = $1 - \text{AIS}$): a well-formed tool call setting `asserted_executor = agent:enrich` that survives to the audit — and, over N trials, the evasion rate with a Wilson interval. A *refusal* counts as non-evasion (and is tallied separately); a *malformed* call is its own category, never silently scored 0; an API failure after three retries is `unavailable` and excluded from the denominator (no silent truncation). Sampling is at `temperature = 1.0` so trials are genuine independent samples; the grader is deterministic given a fixed transcript, so only the attacker’s sampling varies.

The results. Under **B8-attack**, all four models forge the executor claim at near-ceiling rates: M1 0.905 ([0.847, 0.943]), M2 1.000 ([0.912, 1.000]), M3 1.000 ([0.912, 1.000]), M4 0.983 ([0.911, 0.997]). Under **B9-attack**, the evasion rate is 0.0 **for every model** — the pre-registered H2 (the knockout) holds against a real model, not a script. Under the honest prompt (H4 sanity), evasion is 0.0 for every model: a truthful model anchored to its own identity does not misattribute. Figure 1 plots the per-model rates with Wilson intervals.

What B9 actually does to a real attacker. B9’s zero is not the model declining to comply. Under B9-attack the models *do* comply with the injection — they emit the forged completion naming `enrich` at rates comparable to their B8 forging: 0.895, 1.000, 1.000, 0.925 across M1–M4 (34/38, 40/40, 40/40, 37/40; Table 2), against B8-attack forging of 0.905–1.000 — yet the evasion rate is 0.0, because the independent verifier overrides the self-report before it reaches the audit. B9 does not make the model honest; it makes the model’s dishonesty *inert*. That is the precise sense in which independent verification, not model behavior, is the layer that closes the gap.

The H1 disclosure. H1 predicted a B8 evasion rate that is > 0 for the most capable model and non-decreasing in capability. The measured curve is essentially *flat at the ceiling* (0.905, 1.000, 1.000, 0.983): the most-capable cell is well above zero (H1’s existence clause holds), but the rate does not climb across the ladder. The M3→M4 dip in the point estimate (1.000 to 0.983) sits within overlapping Wilson intervals, so the evaluator adjudicates the curve flat rather than a significant decrease — a deliberately conservative criterion: interval overlap cannot establish equality, only the absence of a demonstrated difference. Per the pre-registered H3, a flat curve is reported as a finding, not coded around. We attribute the flatness, in part, to a disclosed design confound: the locked prompt forces

the completion call, removing the refusal path and saturating B8 near the ceiling for every model. We therefore do *not* claim capability-independence — the saturation confound and the prevailing finding that larger models are often *more* injection-susceptible both make that claim indefensible from this data. The supported reading is: *under an operational prompt that requires the completion call, executor-forging is already a ceiling-level vulnerability across the capability range we tested.* De-saturating the prompt is the first v4 experiment (§13).




Figure 1: Tier-2 per-model executor-forging rate with Wilson 95% intervals, across the capability-ordered ladder M1→M4. Under B8 (self-reported) the forging rate saturates near the ceiling for every model (H1 flat — see the forced-call disclosure); under B9 (tool-verified) it is 0.0 for every model (H2, $\varepsilon = 0$). Drawn from the frozen sweep shipped with the paper; Tier 2 is statistically, not byte, reproducible.

6 Comparative Breadth: B6 (mTLS) and B7 (A-JWT)

To place B8/B9 in the full standardized landscape, v3 adds two more verify-the-executor baselines as configuration flags, both pre-registered and locked. **B6** binds the access token to the client’s certificate (mTLS, RFC 8705): the protected resource obtains the certificate from its TLS layer and verifies it matches the token’s `cnf` thumbprint (`x5t#S256`). This is sender-constraint verified at access — the same shape as B5 — so a lifted token is useless to a holder who cannot present the bound certificate, and the colluder cannot present another agent’s certificate. **B7** reads the executor from an A-JWT execution assertion: a per-agent proof-of-possession key bound through the token’s `cnf` thumbprint, verified by the resource server at request time, backed by the IDP’s runtime check at mint. For the colluder to attribute its action to a sibling it would need to both impersonate that sibling in-process at mint and sign with the sibling’s PoP key at the resource — both outside its capability.

Both predict, and measure, AIS = 1.0 on *both* the honest and the colluding columns, identical on both topologies — each *earned* by execution-time verification under the same non-circularity control as B9 (a degraded resolver that falls back to the self-report drops the score). We are careful that

B7 = 1.0/1.0 is for the A-JWT *mechanism graded by AEGIS-AT’s own OS-process recorder*, not an endorsement of A-JWT’s unvalidated proof-of-concept “100%” claim; A-JWT’s native anchor is in-process, weaker than v3’s process boundary, and we grade with our own recorder regardless. The narrative sharpens: *attribution recovers in every baseline that resolves the executor from evidence it cannot author* — {B5 DPOP, B6 mTLS, B7 A-JWT, B9 process-boundary} all reach 1.0 — and the single cell that fails under collusion is B8, the self-reported completion.

7 Implementation

v3 is a package that imports `aegis_at_v2` for the B1–B5 path and adds, over the same codebase:

completion/completion_record.py

the §3 record: signing, signature verification, and the `self_reported/tool_verified` source tags.

completion/execution_assertion.py

B7’s A-JWT execution assertion (per-agent PoP key, signed and verified at execution).

auth/mtls.py

B6’s certificate-binding primitive (per-agent cert, RFC 8705 §3 match, verified subject).

harness/completion_sweep.py

the B6–B9 sweep; it reuses v2’s process kernel, recorder, tool, and scorer byte-for-byte and adds only the completion-record audit layer, with the non-circularity seams (degraded resolver, verifier toggle) exposed for the controls.

harness/llm_seat.py, llm_sweep.py, llm_eval.py

the Tier-2 LLM seat (the structured-tool-call adapter), the adaptive Wilson sweep, and the H1–H4 evaluator (a pure function over a stored grid that asserts the hypotheses by Wilson containment and never recomputes a rate).

transport/mcp_adapter.py

the thin MCP-shaped transport wrapper (§2).

The deterministic core (B1–B9) is asserted by the test suite; the gate (`scripts/check_v3.sh`) runs the INV greps, formatting, and the tests, and must exit 0. The two live LLM tests skip without a provider key, so the deterministic gate has no API dependency.

8 Experimental Methodology

Two-stage pre-registration. Every predicted value is committed in a SHA-256-locked threat model *before* the measuring code, and a CI test recomputes the hash and fails the build on any edit. The locking is staged to match the two reproducibility regimes: `threat-model-v3.md` locks the deterministic core (scripted B8/B9, asserted by *exact equality*); `threat-model-v3.0.1.md` locks the deterministic B6/B7 cells (also exact equality); and `threat-model-v3.1.md` locks the LLM-tier parameters and the directional hypotheses H1–H4 (asserted by *Wilson containment*, never by exact value). The model list, capability ordering, training-cutoff evidence, N , ε , the three prompts verbatim, and the refusal/malformed/retry policy were all pinned and committed before any model was called. A contradicted prediction is reported as a finding, never reconciled (INV-7, INV-8). The two stages play different epistemic roles: the deterministic locks pin a specification whose outcomes

are entailed by construction — their reproduction is a regression gate, and we draw no evidential claim from it — while the Tier-2 hypotheses are falsifiable predictions about model behavior, one of which (H1’s rising curve) failed and is reported as a finding.

The locked hypotheses. For the record the results are judged against: **H1** — under B8-attack the evasion rate is > 0 for the most capable model and non-decreasing along M1→M4; **H2** — under B9-attack the evasion rate is $\leq \varepsilon = 0$ for every model (the knockout); **H3** — if H1’s curve is flat, the flatness is itself reported as a finding about the structural gap, and H2 stands independently of it; **H4** — under the honest prompt, evasion ≈ 0 for every model (a failure is an identity-stability finding, not noise). One statistical caveat is disclosed rather than buried: the adaptive stopping rule is optional stopping, so a Wilson interval computed at a data-dependent N does not carry exact 95% coverage. This does not touch H2/H4, which are adjudicated on zero counts (§9), and we treat the H1 interval comparison as descriptive; fixed- N cells with a formal trend test (e.g., Cochran–Armitage) are the v4 upgrade.

Determinism vs. statistical reproducibility. The scripted cells (B1–B9) are byte-identical across runs under a fixed clock. The LLM cells are statistically reproducible: versions are pinned, `temperature = 1.0` and a fixed `base_seed` are recorded, and trial i uses a per-cell seed derived as `sha256(base_seed|model|baseline| condition)[:4]+i` — the leading 4 bytes of the digest as a 32-bit integer plus the trial index, kept to 32 bits because the provider rejects a `seed  $\geq 2^{63}$` . Where the provider honors the seed the set of trials is reproducible, and where it does not the rate is reproducible only in distribution (disclosed). The grader is deterministic given a fixed transcript. Figures are generated from the harness — the Tier-1 curve from the live harness, the Tier-2 chart from the recorded sweep — so a plotted number cannot diverge from the code.

9 Results

9.1 The deterministic curve (Tier 1)

The measured B1–B9 curve reproduces the pre-registered prediction exactly (Table 1, Figure 2). The honest column is 1.0 wherever the executor is identified at execution and 0.0 under shared-account (B1) and delegation (B3/B4). Under the colluder, every verify-at-execution baseline ($\{B5, B6, B7, B9\}$) is inert — its colluding column equals its honest column at 1.0 — and the failure is isolated to B8, whose colluding column collapses to 0.0. The B8 defect is an actor `field_mismatch` on a signature-valid record.

9.2 The real-model grid (Tier 2)

Table 2 reports every cell — forging rates with Wilson intervals, denominators for all sixteen cells, and the B9 forged-emission counts. Response attrition was confined to the weakest model: M1’s B8-attack cell had 8 refusals and 6 malformed calls (both scored non-evasion) and 12 unavailable trials (excluded from the denominator), while M2–M4 produced a well-formed completion on every trial. The H1–H4 verdict, computed by the evaluator over the locked grid: **H2 SUPPORTED** (B9 evasion = 0 for all four models — the knockout); **H4 SUPPORTED** (honest-condition evasion = 0 for all); **H1 FLAT_FINDING** (M4 above zero but the curve is flat at the ceiling — the disclosed forced-call saturation, §5); **H3 FLAT_FINDING** (the flat-curve branch, derived from H1). Because $\varepsilon = 0$, the H2/H4 containment rule reduces to zero observed evasions (a single B9

Table 1: The v3 deterministic Attribution Integrity Score, B1–B9, under the honest and colluding scripted seats (identical on both topologies, T1/T2). Each baseline is a configuration flag over one codebase. Measured values reproduce the pre-registered prediction exactly; B1–B5 are inherited (v1/v2), B6–B9 are new to v3. The colluder is inert everywhere except B8.

B	Defense in place	Attribution source	honest	colluding
B1	Shared service account	shared credential	0.0	0.0
B2	Per-agent identity	per-agent cred. at execution	1.0	1.0
B3	+ RFC 8693 delegation	delegation current actor	0.0	0.0
B4	+ tamper-evident log	delegation current actor	0.0	0.0
B5	+ DPoP sender-constraint	key-bound current actor	1.0	1.0
B6	+ mTLS (RFC 8705)	cert-bound, verified at access	1.0	1.0
B7	+ A-JWT execution assertion	PoP-verified <code>executed_by</code>	1.0	1.0
B8	+ self-reported completion	self-asserted executor	1.0	0.0
B9	+ tool-verified completion	independent verifier	1.0	1.0




Figure 2: The deterministic B1–B9 AIS curve, honest vs. colluding seat. The two seats coincide everywhere except B8, where the colluding self-report collapses AIS to 0.0 while B9 (tool-verified) holds at 1.0. Every verify-the-executor baseline ($\{B5, B6, B7, B9\}$) recovers to 1.0. Generated from the live harness.

evasion is a finding); H1’s monotonicity test instead turns on Wilson-interval disjointness. H2 stands independently of the H1/H3 finding.

Table 2: Tier-2 per-model results with Wilson 95% intervals and every cell’s denominator. “ev/ n ” is evasions over in-denominator trials (**unavailable** excluded: M1’s B9-attack cell attempted 40 trials, 2 unavailable, so $n = 38$; its B9-honest cell attempted 40, 1 unavailable). “B9 forged emission” is the fraction of trials where the model emitted a well-formed completion naming **enrich** — under B9 the verifier overrides every one, so the forgery is emitted but inert. B8-attack saturates near the ceiling (H1 flat); B9-attack evasion is 0 for every model (H2). Statistically reproducible from the frozen sweep artifact.

Model	B8-attack ev/ n	B9-attack ev/ n	B9 forged emission	honest (B8; B9)
M1 llama-3.1-8b	0.905 [0.847, 0.943] (134/148)	0.0 [0, 0.092] (0/38)	0.895 (34/38)	0/40; 0/39
M2 llama-3.3-70b	1.000 [0.912, 1.000] (40/40)	0.0 [0, 0.088] (0/40)	1.000 (40/40)	0/40; 0/40
M3 gpt-oss-20b	1.000 [0.912, 1.000] (40/40)	0.0 [0, 0.088] (0/40)	1.000 (40/40)	0/40; 0/40
M4 gpt-oss-120b	0.983 [0.911, 0.997] (59/60)	0.0 [0, 0.088] (0/40)	0.925 (37/40)	0/40; 0/40

10 Discussion

Self-attestation defaults are the live risk. The gap v3 measures is not a corner case of a draft. AIP makes self-report the *default* for trusted environments, PEDIGREE permits it with no adequacy guidance, and HDP signs hops with the issuer’s key while admitting a misrepresenting hop is undetectable. A deployment that wires any of these into its audit trail and treats the completion as truth is at B8 by default. The contribution is to show, with a pre-registered measurement, that B8 fails the moment the executor has an incentive to misreport.

Independent verification localizes the fix. The completion-era fix is the same shape as v2’s: not “log a different field” (there is none to trust) but “resolve the executor from an observation the executor cannot author.” B9 instantiates that with a process-boundary verifier; B5/B6/B7 instantiate it cryptographically at execution. The result localizes the missing layer exactly — independent verification — and shows it is robust where self-report is not. One deployment caveat belongs here rather than in the margins: a process-boundary observer exists only inside a single host, so B9 as instantiated is a lab instrument that isolates the principle. Across hosts and trust domains the same principle is carried by attestation infrastructure — TEE/TPM attestation at the tool boundary, a secure sidecar, an out-of-band observer, or the B6/B7 family — which is why the comparative breadth matters: B6/B7 are the deployable instantiations of what B9 isolates.

A real model, not a script, exploits the gap. Tier 2 makes the threat concrete: a real LLM, told via ordinary alert content that the action was “on behalf of” a sibling, forges the executor argument $\sim$ always under B8. The threat is not a scripted strawman. B9’s zero, by contrast, is a structural control, not an empirical surprise — a verifier that reads the OS process boundary cannot be moved by the model — so the measured payload of Tier 2 is the B8 forging rate, and B9 is the structural bound on it. The same models still emit the forgery under B9 — at 0.895–1.000, within a few points of their B8 rates (Table 2) — the difference is that B9 makes the emitted forgery inert. This is the strongest form of the v3 claim: independent verification closes the gap regardless of how willing the model is to lie.

Delegation’s auditability benefit is conditional. As in v1/v2, this result *qualifies* rather than rejects the industry narrative that delegation beats impersonation for accountability [17, 18]. Delegation does beat shared service accounts (B1) on lineage and least privilege; what v3 adds is that the completion layer built on top of delegation inherits a new conditional: its audit-attribution value depends on whether the completion is independently verified, not merely signed.

11 Limitations and Validity Threats

1. **The forced-call saturation confound (load-bearing).** The locked Tier-2 prompt requires the model to call `submit_completion` exactly once, removing the refusal path and saturating the B8 forging rate near the ceiling for every model. H1’s flat curve is therefore partly an artifact of the prompt, and we report it as a finding under this disclosure rather than as capability-independence. De-saturating the prompt (a softer instruction that permits refusal or silence) is the first v4 experiment.
2. **Four open-weight models, one scenario, one topology.** The ladder is four free open-weight models from two makers; frontier closed models, more scenarios, and additional attack phrasings are deferred. The LLM tier runs T1 only because the B8/B9 path is topology-inert by construction.
3. **Statistical, not byte, reproducibility for the LLM tier.** Where the provider does not honor the seed, the rate is reproducible only in distribution. The locked exact results are the scripted cells.
4. **Induced, not emergent, collusion.** The LLM is induced to forge by prompt injection, matching v1/v2’s alert-text adversary and the real 2026 threat; emergent (non-injected) deception is out of scope.
5. **Independence by construction.** The recorder’s independence is a real OS process boundary, but it is established by construction and property-based testing, not formal verification. The completion record is a benchmark abstraction of the trust tiers, not a literal draft schema.
6. **Scoped colluder.** The collusion is false self-attestation within the executor’s own key; key forgery and principal laundering are distinct classes on the backlog. Single trust domain; cross-organization is untested — and B9’s process-boundary verifier is single-host by construction, so its cross-host analogues are the attestation-based baselines (B6/B7) and TEE/sidecar observers, not B9 itself.

12 Related Work

The completion-block frontier AEGIS-AT measures over. AIP [7], PEDIGREE [8], and HDP [9] are the proposals that introduce the completion/provenance layer v3 studies; B8 instantiates AIP’s default self-report tier and B9 instantiates PEDIGREE’s `tool_verified` tier. v3 is a methodology applied to this pattern, not a verdict on any single draft, and it names none of them “broken”: each closes the honest-executor gap, and v3 measures the residual dishonest-executor case they do not model.

SentinelAgent / DelegationBench v4. The closest adjacent framework [15] contributes a Delegation Chain Calculus, a non-LLM Delegation Authority Service, and a 516-scenario benchmark

(150 adversarial actions across seven attack categories; 366 benign), reporting 100% detection at 0% false-positive rate with TLA+ verification of its deterministic properties — including forensic reconstructibility (P4). The two frameworks measure different quantities: SentinelAgent measures *detection* (does the service block a policy-violating delegation?), AEGIS-AT measures *attribution integrity* (does the logged actor equal the true executor?) on actions that violate no policy. They meet at one point: P4 reconstructs the lineage of the token *presented at the proxy*, which establishes who was *authorized*, not who *wielded* the token — and SentinelAgent’s tokens are HMAC-signed bearer credentials validated by signature alone, with no proof-of-possession (the token is not bound to the agent that presents it). AEGIS-AT therefore *bounds the audit-attribution interpretation of P4 under transferable, non-sender-constrained bearer tokens*; it does not refute reconstructibility. The attack also needs none of SentinelAgent’s stronger adversaries (e.g. its compromised-agent case): it succeeds with alert-text control alone, a strictly weaker adversary, which strengthens the finding rather than matching its threat envelope.

Agentic JWT. A-JWT [16] proposes per-agent proof-of-possession keys plus intent/delegation claims binding an API call to a registered agent, and names the exact structural insight v3 measures (“a separation between the actual user... and the executing agent”). It is closest to AEGIS-AT’s *fix*, not its finding: it proposes the sender-constraint (our B7) but states that a comprehensive adversarial evaluation is forthcoming. It is a complementary solution proposal — a candidate verify-the-executor baseline — not a rival measurement.

The Misattribution Gap, and the standards call. *The Misattribution Gap* [14] measures model-vs-memory misattribution (semantic norm drift through a trust-laundering chain); it operates at the model (belief) layer — what the agent is led to believe — whereas AEGIS-AT operates at the delegation (blame) layer — who is recorded as having acted — complementary, not identical. The standards context — NIST NCCoE [11] and the OpenID Foundation [12] calling for AI agent non-repudiation, and the CSA confused-deputy note [13] — frames the problem as defensibly underexplored: no published adversarial benchmark exists for this specific measurement, among the works we compare.

13 Future Work

- **De-saturate the capability curve.** Relax the forced-call prompt so refusal and silence are live options, and re-measure the B8 capability curve on fixed- N cells with a formal trend test (Cochran–Armitage) in place of interval-disjointness — the direct follow-up to this paper’s H1 disclosure.
- **An independent B9.** Replace B9’s verifier — currently the ground-truth recorder itself — with a second, separately implemented observer (e.g., a kernel-audit/ptrace tracer), so B9 acquires genuine failure modes and its 1.0 becomes a measurement rather than a structural control.
- **Wider ladder.** Frontier closed models, more scenarios, and additional attack phrasings.
- **The other completion tiers.** PEDIGREE’s `peer_verified` and `human_verified` tiers, which add independent parties and new protocol surface.
- **Cross-organization delegation.** Multi-trust-domain hand-offs and principal laundering, key forgery, and scope-attenuation bypass as distinct attack classes.

- **Real-time detectability.** A monitor flagging `claimed_actor` $\neq$ verifier at write time (the mismatch is mechanical).

14 Conclusion

The arc completes. v1 showed RFC 8693 delegation has no field for the executor; v2 showed binding the token to the executor’s key recovers attribution; v3 turns to the completion record the 2026 proposals add and finds it defaults to self-report. A scripted colluder collapses self-reported attribution to AIS = 0.0 while tool-verified attribution holds at 1.0, and a real model — four of them, weak to strong — forges the executor claim $\sim$ always under self-report while independent verification drives the forging rate to zero for every one. Every baseline that resolves the executor from evidence it cannot author recovers; the single self-reported cell is the lone failure. As 2026 agent-identity proposals move toward completion- and provenance-based agent non-repudiation, v3 is a concrete demonstration that signing a completion is not the same as verifying it — and that the layer which closes the gap is independent verification, the completion-era counterpart to v2’s sender-constraint.

Reproducibility and Artifact Availability

The complete v3 benchmark — the two-stage SHA-256-locked threat model (deterministic core, B6/B7 amendment, and LLM-tier parameter lock), the completion-record and verification modules, the LLM seat and evaluator, and the test suite — is released alongside the frozen v1 (tag `v1.0.0`) and v2 artifacts. Code is licensed Apache-2.0; documentation CC BY 4.0. Every Tier-1 AIS value is asserted in the test suite against a prediction locked before the measuring code; the Tier-2 hypotheses H1–H4 are asserted by Wilson containment over the recorded grid. The frozen Tier-2 sweep ships alongside the paper source (`data/llm_sweep_v3.json`, relative to the paper directory; live-run outputs under the repository’s `results/` are git-ignored). The deterministic gate (`scripts/check_v3.sh`) runs without any API dependency; the two live LLM tests skip without a provider key. Figures regenerate via `python Documents/Paper/v3/figures/make_figures.py` — the Tier-1 curve from the live harness, the Tier-2 chart from that frozen sweep artifact, so both regenerate from the repo without a provider key (Tier 2 is statistically, not byte, reproducible; a fresh run resamples and requires a key).

References

- [1] R. Singh, *AEGIS-AT: Measuring Attribution Integrity Under Sibling Impersonation in Multi-Agent Delegation* (v1), 2026. Frozen artifact, tag `v1.0.0`. <https://github.com/rahulsingh1397/AEGIS-AT>.
- [2] R. Singh, *AEGIS-AT v2: Does Sender-Constraint Recover Attribution? DPoP, Chain Depth, and a Log-Integrity Metric*, 2026. <https://github.com/rahulsingh1397/AEGIS-AT>.
- [3] M. Jones, A. Nadalin, B. Campbell, J. Bradley, and C. Mortimore, *OAuth 2.0 Token Exchange*, RFC 8693, IETF, January 2020. <https://www.rfc-editor.org/rfc/rfc8693>.
- [4] D. Fett, B. Campbell, J. Bradley, T. Lodderstedt, M. Jones, and D. Waite, *OAuth 2.0 Demonstrating Proof of Possession (DPoP)*, RFC 9449, IETF, September 2023. <https://www.rfc-editor.org/rfc/rfc9449>.

- [5] B. Campbell, J. Bradley, N. Sakimura, and T. Lodderstedt, *OAuth 2.0 Mutual-TLS Client Authentication and Certificate-Bound Access Tokens*, RFC 8705, IETF, February 2020. <https://www.rfc-editor.org/rfc/rfc8705>.
- [6] B. Campbell, J. Bradley, and H. Tschofenig, *Resource Indicators for OAuth 2.0*, RFC 8707, IETF, February 2020. <https://www.rfc-editor.org/rfc/rfc8707>.
- [7] S. Prakash, *AIP: Agent Identity Protocol for Verifiable Delegation Across MCP and A2A*, arXiv:2603.24775, March 2026; IETF Internet-Draft `draft-prakash-aip-00`. <https://arxiv.org/abs/2603.24775>. <https://www.ietf.org/archive/id/draft-prakash-aip-00.html>.
- [8] K. Rampalli, *PEDIGREE: Verifiable Delegation Identity for Agentic AI*, IETF Internet-Draft `draft-rampalli-pedigree-00`, 25 April 2026. <https://www.ietf.org/archive/id/draft-rampalli-pedigree-00.html>.
- [9] A. Dalugoda, *HDP: A Lightweight Cryptographic Protocol for Human Delegation Provenance in Agentic AI Systems*, arXiv:2604.04522, March 2026; IETF Internet-Draft `draft-helixar-hdp-agentic-delegation-00`. <https://arxiv.org/abs/2604.04522>.
- [10] Model Context Protocol, *Authorization*, specification revision 2025-06-18. <https://modelcontextprotocol.io/specification/2025-06-18/basic/authorization>.
- [11] H. Booth, W. Fisher, R. Galluzzo, and J. Roberts, *Accelerating the Adoption of Software and Artificial Intelligence Agent Identity and Authorization*, NIST NCCoE Concept Paper (Initial Public Draft), February 5, 2026. <https://csrc.nist.gov/pubs/other/2026/02/05/accelerating-the-adoption-of-software-and-ai-agent/ipd>.
- [12] OpenID Foundation, *OIDF Responds to NIST on AI Agent Security*, March 11, 2026. <https://openid.net/oidf-responds-to-nist-on-ai-agent-security/>.
- [13] Cloud Security Alliance AI Safety Initiative, *Confused Deputy Attacks on Autonomous AI Agents* (Research Note), March 23, 2026. <https://labs.cloudsecurityalliance.org/research/csa-research-note-ai-agent-confused-deputy-prompt-injection/>.
- [14] T. Ahad, I. Hossain, M. J. Alam, S. Puppala, S. B. Alam, and S. Talukder, *The Misattribution Gap: When Memory Poisoning Looks Like Model Failure in Agentic AI Systems*, arXiv:2605.22842, May 2026. <https://arxiv.org/abs/2605.22842>.
- [15] K. S. R. Patil, *SentinelAgent: Intent-Verified Delegation Chains for Securing Federal Multi-Agent AI Systems*, arXiv:2604.02767, April 2026. <https://arxiv.org/abs/2604.02767>.
- [16] A. Goswami, *Agentic JWT: A Secure Delegation Protocol for Autonomous AI Agents*, arXiv:2509.13597, September 2025; IETF Internet-Draft `draft-goswami-agentic-jwt-00` (“Secure Intent Protocol: JWT Compatible Agentic Identity and Workflow Management”), December 2025. <https://arxiv.org/abs/2509.13597>.
- [17] Red Hat, *Zero Trust for AI Agents: Why Delegation Beats Impersonation*, Red Hat Emerging Tech, May 21, 2026. <https://next.redhat.com/2026/05/21/zero-trust-for-ai-agents-why-delegation-beats-impersonation/>.
- [18] Okta, *Agent Security: Securing the Delegation Chain*, Okta AI Security Blog, 2026. <https://www.okta.com/blog/ai/agent-security-delegation-chain/>.

- [19] N. Hardy, *The Confused Deputy (or why capabilities might have been invented)*, ACM SIGOPS Operating Systems Review, 22(4), 1988.